

Pattern Recognition Report

Orion Mathews
Imperial College London
omm15@ic.ac.uk
01053855

Paul Roever
Imperial College London
pr2214@ic.ac.uk
00948609

Abstract

Face recognition is a complex pattern recognition problem. Principle component analysis (PCA) is a standard technique used in face recognition, however discriminative information between classes may be lost in projection. To maximally separate data from different classes linear discriminant analysis (LDA) is preformed. LDA often provides better face recognition accuracy than PCA. A drawback of LDA is that there must be more training examples than features. Fisher-faces is a method used to solve this problem by initially using PCA to reduce the number of features and then preforming LDA. Face recognition performance can be further improved by fusing together an ensemble of different models.

1. Introduction

A set of faces can be described by a list of features. However, some features may be related and therefore redundant. PCA is used to extract the relevant information in the data set by obtaining the features that show the most data variation. Therefore the data can be represented using fewer features as the unnecessary correlations among the original set of features is removed.

PCA can be used for face recognition using a classifier such as nearest neighbours (NN). However, PCA simply finds the features that are most useful for representing data and may discard discriminating information between the classes. Linear discriminant analysis (LDA) finds the features that best separate data of different classes and usually offers better performance.

Each face image is 46x56 pixels and has been raster-scanned so that it forms a D dimensional column vector (where $D = 46 * 56 = 2576$ and is regarded as the number of features describing each face). The data set X contains 520 images of 52 different people, consequently there are 10 images of each person's face.

2. Cross Validation

Cross validation is used to assess how well the predictive model will generalize to unseen data. This is achieved by splitting the data set into two independent subsets, the training data set and the test data set. The test data remains unknown to the model, allowing the model's accuracy to be tested on new data. Cross validation reveals problems such as models under-fitting or over-fitting the data, due to the models being too simple or complex.

The data set X is partitioned into training data X_{train} and test data X_{test} by defining the number of images per person that will be used for training. The number of training images per person is defined as 7, causing the data set to be split in a ratio of 7:3 training data to test data. Therefore models are trained on 70% of the data (364 images = N_{train}) and tested on the remaining 30% (156 images = N_{test}).

For implementation of data partitioning see lines 19 to 28 in the source code (appendix A).

3. Eigenfaces (PCA)

3.1. PCA applied to training data set (X_{train})

Each face is a D -dimensional vector in the image feature space. A subspace is computed such that the projections of the data points onto the subspace have maximum data variance, this subspace is called the face-space (M dimensional, $M \ll D$). The face-space is spanned by the eigenfaces (the features that characterise variations among face images). Eigenfaces are the principle components of the set of face images, that is, they are the eigenvectors that correspond to the M largest eigenvalues of the face data's covariance matrix S .

$$S = \frac{1}{N_{train}} A * A^T \quad (1)$$

where N_{train} is the number of training images and A is the normalised training data.

Eigenvalues represent the amount of data variation associated with a certain direction. The larger the eigenvalue the

more the data varies in that direction. An eigenvalue of zero means there is no data variation in that direction. Since S is real and symmetric all its eigenvalues are real. The eigenvectors associated with the M largest eigenvalues are chosen as it ensures the chosen directions have maximum data variance. When implementing face recognition with PCA and NN classification, M is chosen so that the sum of the first M eigenvalues is above 99% of the total sum of all the eigenvalues, therefore more than 99% of the information is retained. Given this condition, M is calculated to be 243. See lines 54 to 57 in source code for implementation. Figure 1 shows the value of each eigenvalue.

To map the training faces to the face-space, the mean face of the training data (see figure 2) is subtracted from each face and the result is projected onto the face-space. For the detailed procedure see slides 15 to 20 of the "Face Recognition by Eigenfaces: Subspace, PCA" lecture notes. Figure 3 shows the top 15 eigenfaces (the eigenvectors corresponding to the 15 largest eigenvalues of S).

The number of non-zero eigenvalues is the number of directions in which the data varies and is equivalent to the rank of the matrix S . $A \in \mathbb{R}^{D \times N_{train}}$ and has columns consisting of the normalised training images.

$$rank(A) \leq \min(D, N_{train}) = N_{train} \quad (2)$$

$$rank(A) = rank(A^T A) = rank(AA^T) \quad (3)$$

From equations (1), (2) and (3) it is clear that

$$rank(S) = rank(A) \leq N_{train} = 364 \quad (4)$$

As A is X_{train} normalised, the maximum number of non-zero eigenvalues can be further bounded to

$$rank(S) \leq N_{train} - 1 = 363 \quad (5)$$

Therefore the number of non-zero eigenvalues is the number of training faces -1 ($N_{train} - 1$).

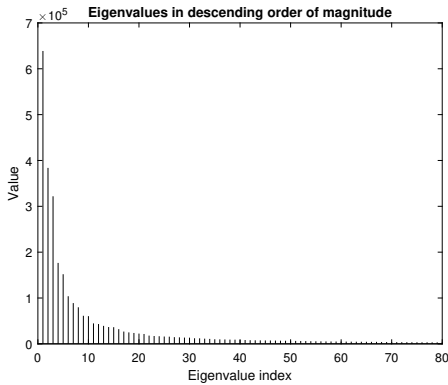


Figure 1. Eigenvalues in descending order




Figure 2. Average training data face

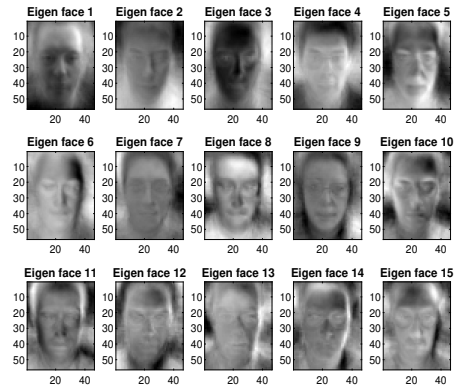


Figure 3. Eigenfaces corresponding to largest 15 eigenvalues.

3.2. Low dimension computation of the face-space

Using (1) to compute S means that $S \in \mathbb{R}^{D \times D}$, as $D = 2576$ this matrix is very large, hence eigenvalue-eigenvector decomposition of S is computationally expensive. Therefore, we consider the eigenvalue-eigenvector decomposition of

$$\frac{1}{N_{train}} A^T A \quad (6)$$

$$\left(\frac{1}{N_{train}} A^T A \right) * \mathbf{V}_i = \lambda_i \mathbf{V}_i$$

Multiplying by A on the left:

$$\left(\frac{1}{N_{train}} A A^T A \right) * \mathbf{V}_i = \lambda_i A * \mathbf{V}_i$$

Using equation (1) and applying a simple substitution:

$$S A * \mathbf{V}_i = \lambda_i A * \mathbf{V}_i$$

$$S \mathbf{U}_i = \lambda_i \mathbf{U}_i \quad (7)$$

From equation (7) and identity (3) the non-zero eigenvalues of AA^T and $A^T A$ are identical. The corresponding eigenvectors can be obtained as $\mathbf{U}_i = A * \mathbf{V}_i$.

Using (6) to compute the data co-variance matrix results in $S \in \mathbb{R}^{N_{train} \times N_{train}}$; as $N_{train} = 364 \ll 2576$, the matrix is much smaller than if S is computed using (1). The benefits of having a smaller co-variance matrix is that the time taken to compute the eigenvectors and eigenvalues is much less. For $S \in \mathbb{R}^{D \times D}$ the time taken for eigenvalue-eigenvector decomposition is approximately 12.59 seconds compared to around 0.090 seconds when $S \in \mathbb{R}^{N_{train} \times N_{train}}$. A possible shortcoming of this method is that the reconstruction error appears to be larger when using the low dimensional S . Overall the much faster computational speed outweighs this drawback and all further analysis will be carried out with the low-dimensional data co-variance matrix.

3.3. Face reconstruction

Each face can be reconstructed by a linear combination of the eigenfaces. Perfect reconstruction of the original data set can theoretically be achieved if all eigenvectors are used. However using all eigenfaces would be computationally expensive and as shown in table 1 an average face reconstruction accuracy of 91.4% can be achieved with only 100 eigenfaces. As expected the more eigenfaces that are used in reconstruction the higher the average reconstruction accuracy, this is illustrated in figure 4 as well as in table 1.

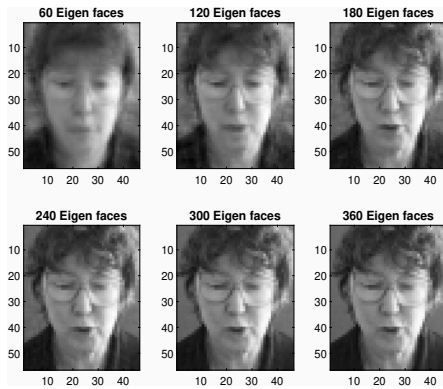


Figure 4. Face reconstruction using different number of eigenfaces

Number of eigenfaces used in reconstruction (M)	Average reconstruction accuracy (%)
40	87.1
80	90.4
120	92.3
160	93.8
200	95.1
240	96.2
280	97.3

Table 1. Table of eigenfaces used in reconstruction and corresponding accuracies

3.4. Nearest Neighbor (NN) classification

Face recognition using nearest neighbor classification is preformed by computing the face-space with all the training data, projecting the new image onto the face-space and classifying the face by comparing its position in the face-space with the positions of known individuals. The new image is assigned to the same class as the training image that is closest to it in the face-space.

As the number of eigenfaces learned increases so does the prediction accuracy (shown in figure 5). Initially the prediction accuracy increases rapidly but plateaus after reaching approximately 63%. After 100 learned eigenfaces there is very little increase in prediction accuracy. With 243 learned eigenfaces, recognition is 63.5% accurate. Figure 7 shows the confusion matrix for NN classification using 100 eigenfaces.




Figure 5. NN classification accuracy against number of eigenfaces learned.



Figure 6. Example of failure case. PCA with NN classification miss-classified the individual on the left. The image on the right is the image that was classed as most similar, though it is a completely different individual.

3.5. Alternative classification method

Alternatively, a face-space can be computed for each person. The new image would then be projected onto each

of the face-spaces and the reconstruction error can be measured. The assigned class would be the one with least reconstruction error.

Face recognition using this method is 74.4% accurate where as the maximum accuracy for NN classification was 61.9%; both were averaged over five repetitions with differently shuffled data. The difference is that in this alternative method similarity to all images is taken into account instead of classifying based on the single most similar image. For comparison the confusion matrix for this alternate method of classification is included in appendix C.




Figure 7. PCA with NN classification confusion matrix (M=100). For full figure see appendix C.

4. Combined Generative and Discriminative Subspace Learning

PCA finds the directions that maximise the projected data variance. These directions are the columns of the projection matrix W_{PCA} . W_{PCA} is found by solving the PCA objective function given in equation (8), where W is a matrix of projection directions and S is the data covariance matrix.

$$W_{PCA} = \arg \max_W |W^T * S * W| \quad (8)$$

with $W^T * W = 1$

LDA finds the directions that maximise the separation of data from different classes. These directions are the columns of the projection matrix W_{LDA} . W_{LDA} is found by solving equation (9), i.e. maximising the between class scatter while minimizing the within class scatter. S_B is the between class scatter matrix and S_W is the within class scatter matrix.

$$W_{LDA} = \arg \max_W \frac{|W^T * S_B * W|}{|W^T * S_W * W|} \quad (9)$$

S_W is invertible if $N_{train} > D$, if this is the case then W_{LDA} is found to be made up of the eigenvectors of $S_W^{-1} S_B$ (see slides 8 to 11 of "Discriminant Analysis: Fisherfaces" lecture notes for proof). If $N_{train} < D$ then S_W is singular and the number of dimensions must be reduced. Therefore PCA is preformed and the data is projected to a lower-dimensioned subspace where $N_{train} > D$ and S_W is invertible. This procedure results in objective function (10).

$$W_{LDA} = \arg \max_W \frac{|W^T * W_{PCA}^T * S_B * W_{PCA} * W|}{|W^T * W_{PCA}^T * S_W * W_{PCA} * W|} \quad (10)$$

This is equivalent to maximising the numerator while keeping the denominator constant, giving equation (11).

$$W_{LDA} = \arg \max_W |W^T * W_{PCA}^T * S_B * W_{PCA} * W|$$

with $W^T * W_{PCA}^T * S_W * W_{PCA} * W = k$ (11)

W_{LDA} can be found using the method of Lagrange multipliers as demonstrated below.

$$L = (W^T * W_{PCA}^T * S_B * W_{PCA} * W) + \lambda(k - W^T * W_{PCA}^T * S_W * W_{PCA} * W) \quad (12)$$

Setting gradient with respect to W to zero gives:

$$2W_{PCA}^T * (S_B - \lambda S_W) * W_{PCA} * W = 0 \quad (13)$$

Expansion and factorisation yields

$$(W_{PCA}^T S_W W_{PCA})^{-1} * (W_{PCA}^T S_B W_{PCA}) W = \lambda W \quad (14)$$

From (14), it is clear that W and λ are the eigenvectors and eigenvalues of $(W_{PCA}^T S_W W_{PCA})^{-1} * (W_{PCA}^T S_B W_{PCA})$. The combined PCA-LDA subspace is therefore given by

$$W_{opt} = W_{PCA} * W_{LDA} \quad (15)$$

5. PCA-LDA Face Recognition with NN classification

To perform LDA the within class scatter matrix S_W must be invertible, as described in section 4. Therefore PCA is used to reduce the dimensions of the training data set so that $N_{train} > D$. The training data is then projected onto the low dimension subspace and LDA is preformed on the projected data. See slides 14 to 27 of the "Discriminant Analysis: Fisherfaces" lecture notes for the procedure and lines 38 to 77 in the source code for the implementation (appendix B).

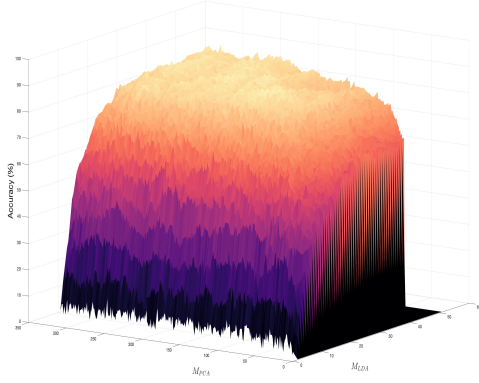


Figure 8. Accuracy heat-map generated by varying the number of eigenvectors used in PCA and LDA. For full image and a heatmap see appendix C.




Figure 9. Confusion matrix for PCA-LDA face recognition with NN classifier. M_{PCA} is 312 and M_{LDA} is 51. For full image see appendix C.

Figure 8 displays the relationship between the face recognition accuracy for every combination of the number eigenvectors used in PCA (M_{PCA}) and LDA (M_{LDA}). If M_{PCA} and M_{LDA} are increased simultaneously the recognition accuracy also increases. The initial accuracy increase is rapid as M_{PCA} increase from 1 to 100 and M_{LDA} increases from 1 to 17 but plateaus after an accuracy of approximately 92%. If M_{PCA} is kept constant the recognition accuracy increases as M_{LDA} increases. The opposite is also true but to a lesser extent. M_{PCA} must always be greater than M_{LDA} to ensure that the within class scatter matrix S_W is invertible.

The rank of the within class scatter matrix (S_W) is the total number of images in the training data (N_{train}) minus the number of classes (C), i.e. $rank(S_w) = N_{train} - C$.




Figure 10. Top 15 Fisher-faces spanning the combined PCA-LDA subspace.

The between class scatter matrix S_B has rank $C - 1$, this is because S_B is a combination of C feature vectors and C feature vectors will define a subspace that has $C - 1$ dimensions or less. Using identity (16) it becomes clear that LDA can reduce dimensionality to at most $C - 1$, since $rank(S_w^{-1}S_b) \leq \min(rank(S_w), rank(S_b)) = rank(S_b) = C - 1$.

$$rank(AB) \leq \min(rank(A), rank(B)) \quad (16)$$

The confusion matrix for face recognition using PCA-LDA with a NN classifier is shown in figure 10. The model is successful at predicting the majority of classes (2 out of 3 test images are predicted correctly for 46 individuals). The average accuracy is 78.8%. The most miss-classified individuals are 1, 19 and 44. In these three cases all three test images were miss-classified. Figures C and 11 are examples of individuals that have been correctly classified and miss-classified respectively.

For a fixed value of M_{PCA} , PCA-LDA with NN classification gives a higher face recognition accuracy than just PCA with NN classification if M_{LDA} is above a certain value. Figure 12 shows how the accuracy of PCA-LDA changes as a function of M_{LDA} when M_{PCA} is fixed at 243. In general, for a given value of M_{PCA} the face recognition accuracy can be improved by using LDA.

6. PCA-LDA Ensemble learning

The performance of the PCA-LDA model can be improved by fusing an ensemble of different models, each one of which has varying parameters or data subsets. This prevents discriminative information from being discarded and over-fitting of the model. The ensemble of randomly different models is created by training each model on random



Figure 11. Example of failure case. PCA-LDA with NN classification miss-classified the individual on the left. The image on the right is the image that was classed as most similar, though it is a different individual. See appendix C for a success case.




Figure 12. Accuracy of PCA and PCA-LDA when M_{PCA} is 312 and M_{LDA} is varied

subsets of the training data and randomising the model parameters.

One approach, known as bagging (discussed in Wang & Tang, 2006), involves splitting the training data into multiple “bags”, each of which contains a subset of the overall training data. Figure 13 explores how many bags to use and how many samples each bag should contain. Although $N_{bags} = 80$ consistently displays slightly higher accuracy than $N_{bags} = 20$, it was much more computationally expensive to use 80 bags. Consequently, 20 bags were used, each containing 250 samples, corresponding to the highest expected accuracy for that number of bags.

Another approach, subspace randomisation[1], is performed by fixing a certain number of the PCA-LDA covariance matrix’s eigenvectors (those with the largest eigenvalues). Then, a random subset of the remaining ones is chosen. W_{PCA} , W_{LDA} , and W_{opt} are then calculated from just the selected eigenvectors. The optimal combination of eigenvectors was found to be 130 total with the 45 largest

remaining fixed and the remaining ones randomly chosen.

Table 2 shows the various accuracies achieved with the different ensemble methods. The baseline Fisherface performance is given at the top. Both ensemble methods, when combined with majority voting (whichever class is predicted the most), already perform better than the baseline. Note that the best-performing bag and subspace performed worse than the combinations of bags and subspaces, respectively. This implies that the random sampling approach works; each LDA model is tuned to unique aspects of the overall feature space, and combining them improves on their individual performance. Instead of combining the individual models using majority voting, the sum rule can be used to improve accuracy slightly further. See appendix C for PCA-LDA ensemble confusion matrix.

Face recognition and classification is a complex problem. PCA + NN is a good starting point for the task. Combining PCA with LDA gives the Fisherfaces algorithm, which performs much better than “pure” PCA. An ensemble approach randomises certain parameters of multiple discriminant models, allowing certain models to specialise to certain aspects of the recognition problem. Combining these results with a committee machine yields more modest, but still substantial, performance gains.




Figure 13. Accuracy of bagging (by majority voting) when $N_{bags} = 20, 80$ as a function of $N_{samples}$ per bag

	Average accuracy (%) over 20 trials
Baseline Fisherface	85.1
Bagging (Majority Voting)	87.7
Random Subspaces (Majority Voting)	89.4
Best-performing Bag	82.2
Best-performing Subspace	85.1
Combined subspaces & bags (Maj. Vot.)	89.9
Combined subspaces & bags (Sum Rule)	90.8

Table 2. Table of average accuracies for different types of ensemble PCA-LDA.

References

- [1] Wang, X. & Tang, X., 2006. Random Sampling for Subspace Face Recognition. International Journal of Computer Vision, 70(1), pp.91–104.

A. Source code Q1

```
1 clear
2 close all
3 clc
4
5 load('face.mat');
6 rng(1)
7
8 %% Parameters
9
10 [D_features, N_faces] = size(X);
11     % N_faces is number of faces (520)
12     % D_features is number of features/dimensions
13     % (2576)
14 N_faces_per_person = 10;
15     % sets the number of faces per
16     % person
17 N_people = N_faces / N_faces_per_person;
18 train_per_person = 7;
19     % sets the split between number of
20     % training faces and number of test faces (per
21     % person)
22 test_per_person = N_faces_per_person -
23     train_per_person;
24 M_max = train_per_person * N_people;
25     % M is number of eigen values used
26     % (M_max is largest possible M = number of
27     % training faces)
28 img_width = 46;
29 img_height = 56;
30
31 %% Splitting data set (for cross-validation)
32
33 train_split = [ones(1, train_per_person), zeros
34     (1, test_per_person)];
35 train_split = train_split(randperm(
36     N_faces_per_person)); % shuffles the
37     % train and test images
38 train_indices = logical(repmat(train_split, 1,
39     N_people)); % logical array that chooses
40     % which columns should be in the training data
41
42 l_train = l(:, train_indices);
43 l_test = l(:, ~train_indices);
44 X_train = X(:, train_indices);
45 X_test = X(:, ~train_indices);
46
47 %% PCA (S = 1/N * A * A')
48
49 X_train_avg = mean(X_train, 2); % compute average
50     % face vector (the mean value of each row/
51     % feature)
52 A = X_train - X_train_avg; % subtract
53     % average face vector from each face/column
54 A_test = X_test - X_train_avg;
55
56 tic
```

```
38 S = (1/N_faces)*A*(A'); % Computing
39     % covariance matrix
40 [U, Eval] = eigs(S, N_faces); % Computing M
41     % eigen-vectors and eigen-values of covariance
42     % matrix
43
44 toc
45
46 Eval = real(sum(Eval, 1));
47
48 %% Plot Eigenvalues
49
50 figure
51 Eval_plot = stem(Eval); % shows the value
52     % of each eigen value. largest eigen values
53     % correspond to best eigen vectors as data has
54     % most variance in these directions
55 set(Eval_plot, 'Marker', 'none');
56 title('Value of each eigen value in descending
57     % order of magnitude')
58 xlabel('Eigenvalue index')
59 ylabel('Value')
60 print -deps Eval_plot
61
62 %% Choosing M
63
64 M = find(0.99*sum(Eval) < cumsum(Eval)); %
65     % finds all indices where the cumulative sum
66     % of the eigen values are greater than 99% of
67     % the total sum of the eigen values
68 M = M(1); % M
69     % is the first element that meets this criteria
70
71 %% PCA_low (S = 1/N * A' * A)
72
73 tic
74 S_low = (1/N_faces)*(A')*A; % low dimensional
75     % computation of the eigenspace
76 [V_low, Eval_low] = eigs(S_low, M_max);
77 U_low = A*V_low;
78 U_low = normc(U_low);
79
80 toc
81
82 Eval_low = real(sum(Eval_low, 1));
83
84 %% Reconstruction while varying number of eigen
85     % faces
86
87 %W_project = U'*A; % high dimension
88     % reconstruction (less error)
89 %X_train_rec = X_train_avg + U*W_project;
90
91 X_face_rec_mat = zeros(D_features, 18);
92 rec_accuracy_mat = zeros(18, M_max);
93 Eval_index = zeros(1, 18);
94 j = 1;
95 for i = 20:20:M_max
96     W_project_low = (U_low(:, 1:i))'*A;
97     X_train_rec_low = X_train_avg + U_low(:, 1:i)
98         *W_project_low;
99
100     X_face_rec_mat(:, j)=X_train_rec_low(:, 1);
101
102     X_train_rec_error = abs(X_train -
103         X_train_rec_low);
104     rec_accuracy = 100 - ((sum(X_train_rec_error)
105         ./ sum(X_train)) * 100);
106     rec_accuracy_mat(j,:) = rec_accuracy;
```



```

88     Eval_index(j) = i;
89     j = j + 1;
90 end
91
92 rec_accuracy_avg = mean(rec_accuracy_mat, 2);
93
94 %% plot fce renonstructions
95
96 figure
97 h = zeros(1, 6);
98 for i = 1:6
99     h(i) = subplot(2, 3, i);
100    img_1 = reshape(X_face_rec_mat(:, 3*i), [
101        img_height, img_width]);
102    image(img_1, 'Parent', h(i));
103    title([num2str(i*60) ' Eigen faces']);
104 end
105 colormap(gray(255));
106 print -deps face_recs
107
108 %% plot reconstruction accuracy
109
110 figure
111 stem(Eval_index, rec_accuracy_avg);
112 title('Reconstruction accuracy as against the
113     number of Eigenfaces used')
114 xlabel('Number of Eigenfaces used in
115     reconstruction')
116 ylabel('Reconstruction accuracy (%)')
117 print -deps Rec_acc
118
119 %% KNN classification using PCA
120
121 predicted_class_matrix = zeros(M_max,
122     test_per_person * N_people); % stores
123     predicted classes for each face
124 accuracy_matrix = zeros(1, M_max);
125     % stores accuracies for
126     each number of eigen values used
127
128 for i = 1:M_max
129
130     X_train_proj = (A')*U_low(:, 1:i); %
131     Projecting normalised faces onto the face-
132     space. Rows are the projections of normalised
133     faces onto the eigenface (row 1 is
134     projection of face 1)
135     X_test_proj = (A_test')*U_low(:, 1:i);
136
137     Idx = knnsearch(X_train_proj, X_test_proj);
138     % Perform a knnsearch between X_train_proj
139     and X_test_proj to find indices of nearest
140     neighbor and puts in coloumn vector
141
142     predicted_class = 1_train(Idx);
143     prediction_results = predicted_class ==
144     1_test;
145     prediction_accuracy = (sum(prediction_results)
146     *100)/(test_per_person * N_people);
147
148     predicted_class_matrix(i, :) =
149     predicted_class;
150     accuracy_matrix(i) = prediction_accuracy;
151 end
152
153 %X_train_rec = X_train_avg + U_low*X_train_proj';

```

```

154 %Error_rec = vecnorm(X_train - X_train_rec);
155
156 %% Plot average training face
157
158 figure
159 img_1 = reshape(X_train_avg, [img_height,
160     img_width]); % takes the average face vector
161     (D x 1) and changes it to an image matrix (W
162     x H)
163
164 image(img_1);
165 colormap(gray(255));
166 print -deps PCA_avg_face
167
168 %%
169
170 figure
171 plot(accuracy_matrix); % shows can stop at
172     around 100 eigen values
173 title('Classification accuracy against number of
174     eigen values used')
175 xlabel('Eigen values used (M)')
176 ylabel('Accuracy (%)')
177
178 figure
179 %confusionchart(1_test, predicted_class_matrix
180     (100, :)); % selects the results obtained
181     using 100 eigen vectors and creates a
182     confusion chart from them
183
184 %% Classification using reconstruction
185
186 Error_rec = zeros(N_people, test_per_person *
187     N_people);
188
189 for i = 1:N_people
190     index = (i - 1) * train_per_person;
191
192     class_train = X_train(:, index + 1:(index +
193     train_per_person)); % takes each class of
194     person
195
196     class_avg = mean(class_train, 2);
197     class_normalized = class_train - class_avg;
198     S_class = class_normalized'*class_normalized;
199
200     [V_class, Eval_class] = eigs(S_class,
201     train_per_person - 1);
202     U_class = class_normalized*V_class;
203     U_class = normc(U_class);
204
205     test_class_norm = X_test - class_avg;
206     test_class_projection = U_class'*
207     test_class_norm;
208
209     X_test_rec = class_avg + U_class*
210     test_class_projection;
211     test_difference = X_test - X_test_rec;
212
213     Error_rec(i, :) = vecnorm(test_difference);
214 end
215
216 [Mins, I] = min(Error_rec, [], 1);
217
218 figure
219 %confusionchart(1_test, I);
220
221 results_reconst = 1_test == I;

```



```

191 accuracy_reconst = sum(results_reconst)/length(
192     l_test)*100;
193
194
195 %% Plot Error case
196
197 img_width = 46;
198 img_height = 56;
199
200 figure
201 subplot(1, 2, 1);
202 img_1 = mat2gray(reshape(X_test(:,31), [
203     img_height, img_width])); % takes the average
204     face vector (D x 1) and changes it to an
205     image matrix (W x H)
206 imshow(img_1, 'InitialMagnification', 'fit');
207 title('Test image 31 from person 11');
208
209 subplot(1, 2, 2);
210 img_1 = mat2gray(reshape(X_train(:,123), [
211     img_height, img_width])); % takes the average
212     face vector (D x 1) and changes it to an
213     image matrix (W x H)
214 imshow(img_1, 'InitialMagnification', 'fit');
215 title('Training image 123 from person 18');
216
217 colormap(gray(255));
218 print -depsc PCA_NN_Error

```

B. Source code Q3 fisherfaces

```

1 clear
2 close all
3 clc
4
5 %% setup
6 load('face.mat');
7 rng(1)
8
9 % dimensions
10 width = 46;
11 height = 56;
12
13 % set some Ns
14 N = size(X, 2);
15 N_faces_per_person = 10;
16 N_people = N / N_faces_per_person;
17 N_features = size(X, 1);
18
19 % generate train/test split
20 train = 7;
21 test = N_faces_per_person - train;
22
23 % create logical (boolean) indices to be used for
    splitting 1, X
24 train_split = [ones(1,train), zeros(1,test)];
25
26 % shuffle train/test images
27 train_split = train_split(randperm(
    N_faces_per_person));
28
29 train_indices = logical(repmat(train_split, [1,
    N_people]));
30
31 % split dataset, train is the ones in the indices
    , test is the inverse
32 l_train = 1(:, train_indices);
33 l_test = 1(:, ~train_indices);
34
35 X_train = X(:, train_indices);
36 X_test = X(:, ~train_indices);
37
38 mean_train_image = mean(X_train, 2);
39 mean_class_images = zeros(N_features, N_people);
40
41 %% calculate S.W, S.B
42 S.W = zeros(N_people, N_features, N_features);
43
44 % create a mean image for each person
45 for i = 1:N_people
46     index = (i - 1) * train;
47     current_train = X_train(:, index + 1:(index +
        train));
48
49     mean_class_image = mean(current_train, 2);
50     mean_class_images(:, i) = mean_class_image;
51
52     diffed_class_image = current_train -
        mean_class_image;
53
54     % form within-class scatter matrix
55     S.W(i, :, :) = diffed_class_image *
        diffed_class_image';
56 end
57
58 S.W = reshape(sum(S.W, 1), [N_features,

```

```

    N_features]);
59 diffed_class_mean_images = mean_class_images -
    mean_train_image;
60
61 S_B = (diffed_class_mean_images)*(
    diffed_class_mean_images)';
62
63 S_T = S_B + S.W;
64
65 M_pca = 312;
66 % M_pca = rank(S.W);
67
68 [W_pca, D_pca] = eigs(S_T, M_pca);
69
70 intermediate = (W_pca' * S.W * W_pca) \ (W_pca' *
    S_B * W_pca);
71
72 accuracy_mldas = zeros(1, 51);
73
74 for M_lda = 1:51
75     % M_lda = rank(S_B);
76
77     [W_lda, D_lda] = eigs(intermediate, M_lda);
78
79     W_opt = W_pca * W_lda;
80
81     %% testing W_opt
82     proj_train = (X_train - mean_train_image)' *
        W_opt;
83     proj_test = (X_test - mean_train_image)' *
        W_opt;
84
85     Idx = knnsearch(proj_train, proj_test);
86
87     % use 1 to determine if we got the right
        person, results(n) = 0 if we did
88     predicted_class = l_train(Idx);
89     results = abs(l_test - predicted_class);
90
91     % every time we get it wrong, make it a
        positive number, then turn it
92     % into a logical array, invert and sum it to
        get the number of correct
93     % test examples, divide by the total number
        of tests and turn into a
94     % percentage
95     accuracy = sum(~logical(results))/length(
        l_test)*100;
96     accuracy_mldas(M_lda) = accuracy;
97 end
98
99 accuracy_pure_pca = 63.461538461538460;
100
101 plot(accuracy_mldas)
102 hold on
103 plot(ones(1, 51)* accuracy_pure_pca)
104 xlabel('$M_{LDA}$', 'Interpreter', 'latex')
105 ylabel('Accuracy (%)')
106 legend({'Fisherfaces with varying $M_{LDA}$', '
    Pure PCA'}, 'Interpreter', 'latex', 'Location
    ', 'best')
107
108 %%
109 % %% Plot Confusion Chart
110 %
111 % figure
112 % confusionchart(l_test, predicted_class);

```

```

113 %% print('./figures/q3-confusionmatrix_PCA-LDA.
    png', '-dpng')
114 %
115 %%% Plot Error case
116 %
117 % img_width = 46;
118 % img_height = 56;
119 %
120 % figure
121 % subplot(1, 2, 1);
122 % img_1 = mat2gray(reshape(X_test(:,3), [
    img_height, img_width])); % takes the average
    face vector (D x 1) and changes it to an
    image matrix (W x H)
123 % imshow(img_1, 'InitialMagnification', 'fit');
124 % title('Test image 3 from person 1');
125 %
126 % subplot(1, 2, 2);
127 % img_1 = mat2gray(reshape(X_train(:,111), [
    img_height, img_width])); % takes the average
    face vector (D x 1) and changes it to an
    image matrix (W x H)
128 % imshow(img_1, 'InitialMagnification', 'fit');
129 % title('Training image 111 from person 16');
130 %
131 % colormap(gray(255));
132 %% print('./figures/q3-error-case_PCA-LDA.eps',
    '-depsec','-tiff')
133
134 %%% Plot Success case
135
136 img_width = 46;
137 img_height = 56;
138
139 figure
140 subplot(1, 2, 1);
141 img_1 = mat2gray(reshape(X_test(:,143), [
    img_height, img_width])); % takes the average
    face vector (D x 1) and changes it to an
    image matrix (W x H)
142 imshow(img_1, 'InitialMagnification', 'fit');
143 title('Test image 143 from person 48');
144
145 subplot(1, 2, 2);
146 img_1 = mat2gray(reshape(X_train(:,336), [
    img_height, img_width])); % takes the average
    face vector (D x 1) and changes it to an
    image matrix (W x H)
147 imshow(img_1, 'InitialMagnification', 'fit');
148 title('Training image 336 from person 48');
149
150 colormap(gray(255));
151 print('./figures/q3-success-case_PCA-LDA.png', '-
    dpng')
152
153 %%% Plot 15 Biggest Fisherfaces
154 figure
155 h = zeros(1, 15);
156 for i = 1:15
157     h(i) = subplot(3, 5, i);
158     img_1 = mat2gray(reshape(W_opt(:,i), [
    img_height, img_width]));
159     imshow(img_1, 'InitialMagnification', 'fit');
160     title(['Fisherface ' num2str(i)]);
161 end
162 colormap(gray(255));
163 print('./figures/q3-fisherfaces.png', '-dpng')

```

C. Tables & Figures


Table of eigenfaces used in reconstruction and corresponding accuracies

Number of eigenfaces used in reconstruction (M)	Average reconstruction accuracy (%)
20	83.7
40	87.1
60	89.0
80	90.4
100	91.4
120	92.3
140	93.1
160	93.8
180	94.5
200	95.1
220	95.7
240	96.2
260	96.8
280	97.3


Average accuracies as a function of the number of subspaces used in subspace randomisation

$N_{subspaces}$	Average accuracy (%) over 20 trials
5	86.2
10	84.2
15	90.3
20	86.8
25	88.2
30	87.4
35	86.9
40	86.9
45	88.7
50	89.0
55	90.5
60	88.3
65	85.6
70	89.2
75	89.2
80	86.0
85	89.4
90	90.8
95	88.7
100	87.1


Confusion matrix for PCA using NN classification




**Confusion matrix for PCA using alternative method classification
(i.e. using reconstruction errors)**




Accuracy heat-map generated by varying the number of eigenvectors used in PCA and LDA



Accuracy heat-map generated by varying the number of eigenvectors used in PCA and LDA



Confusion matrix for PCA-LDA face recognition with NN classifier. M_{PCA} is 312 and M_{LDA} is 51.



Example of success case for PCA-LDA with NN classification.


Test image 143 from person 48



Training image 336 from person 48



PCA-LDA ensemble confusion matrix



PCA-LDA ensemble confusion matrix

